

CSC 453

Operating Systems

This class

Who am I?

- <https://pschmitt.net>
- Academics
 - PhD at UCSB
 - Researcher (summers) at ICSI
 - Postdoc then research faculty at Princeton
 - Some other stops... USC, Stanford, UH
- Privacy startup co-founder & CEO, measurement startup COO
- Main focuses
 - Integrating privacy and security into systems that we all use
 - ML for network management
 - Connectivity / privacy in challenging environments



This class

- Too much to cover in one term: survey of important topics
- Course Canvas: <https://canvas.calpoly.edu/courses/168889>
- Syllabus: <https://schmittpaul.github.io/CSC453W26/syllabus.pdf>
- Waitlists...
- What OSes do you all run?

Who are you?

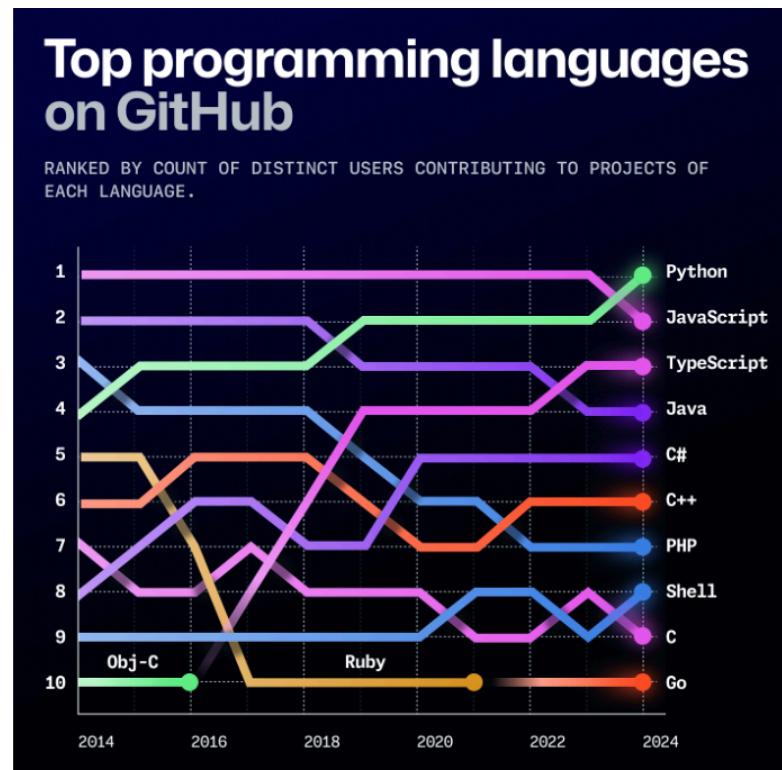
- Name / Major / Year
- Where do you identify as your hometown?

How to succeed in this class

- Turn things in **on time**
- Attend class
 - ... on time. Lectures will start at 10 after the hour
- Ask questions
 - I'm very flexible about how much we cover this semester
 - I would rather teach less and have everyone understand it
 - Our back-and-forth during class is the one of the few indicators I have of how much you are absorbing
- Talk to me if you are struggling

What are we doing here?

- How many of you have participated in OS development?
- How many of you regularly program in languages that use operating system abstractions directly?
- And C is a decreasingly popular language!
- So why study operating systems?
Why is this class even offered?
Why is it required?



Why take this class

- You are required to do so in order to graduate
- **Reality**: this is how computers really work, and as a computer scientist or engineer you should know how computers really work
- **Ubiquity**: operating systems are everywhere and you are likely to eventually encounter them or their limitations
- **Beauty**: operating systems are examples of mature solutions to difficult design and engineering problems. Studying them will improve your ability to design and implement abstractions

Course progression

- Introduction to operating system abstractions and structures
- Abstracting and multiplexing:
 - CPU: interrupts, context, threads, processes, processor scheduling, thread synchronization, deadlocks
 - Memory: memory layout, address translation, paging and segmentation, address spaces, translation caching, page fault handling, page eviction, swapping
 - File systems and storage: disk scheduling, on-disk layout, files, buffer cache, crash and recovery
 - Security (time permitting)

Background

Questions to consider

- What is an operating system?
- What do they do?
- How did OSes evolve?
- How do the requirements differ for different OSes?

What is an operating system?

OS history - how did we get here?

- Started out as **libraries** to provide common functionality across programs
 - See any issues with this approach?
- Later, evolved from **procedure call** to **system call**
 - What's the difference?
- Evolved from running a single program to **multiple processes** concurrently
 - New issues to solve?



Types of operating systems

- Desktop
- Time share/Mainframe
- Mobile
- Web
- Real-time
- Embedded
- Virtual Machines

TL;dr

OSes continue to evolve as environments (i.e., constraints) change

What isn't clear?

Comments? Thoughts?

OS abstractions

Questions to consider

- What are the main abstractions OSes provide?
- What are the abstraction challenges?

Abstractions

- Abstractions *simplify application design* by:
 - hiding undesirable properties,
 - adding new capabilities, and
 - organizing information
- Abstractions provide an **interface** to application programmers that separates **policy**—what the interface commits to accomplishing—from **mechanism**—how the interface is implemented.

What are the abstractions?

- CPUs
 - Processes, threads
- Memory
 - Address space
- Disk
 - Files

Example OS abstraction: file systems

- What **undesirable properties** do file systems hide?
 - Disks are slow!
 - Chunks of storage are actually distributed all over the disk
 - Disk storage may fail!
- What **new capabilities** do files add?
 - Growth and shrinking
 - Organization into directories, searchability
- What **information** do files help organize?
 - Ownership and permissions
 - Access time, modification time, type, etc.

Abstraction tradeoffs - discussion

- Identify undesirable properties hidden by, new capabilities added, and info organization provided with these abstractions:
 - Process / threads
 - Address space

Abstraction pros / cons

- Advantages of OS providing abstractions?
 - Allow applications to reuse common facilities
 - Make different devices look the same
 - Provide higher-level or more useful functionality
- Challenges?
 - What are the correct abstractions?
 - How much should be exposed?

OS design requirements - what do we need?

- Convenience, abstraction of hardware resources for user programs
- Efficiency of usage of CPU, memory, etc.
- Isolation between multiple processes
- Reliability, the OS must not fail
- Other:
 - Security
 - Mobility

What isn't clear?

Comments? Thoughts?

Resource management

Questions to consider

- How does the OS manage access to resources?

OS as a resource manager

- Another view: **resource manager** - shares resources “well”
- **Advantages** of the OS managing resources?
 - Protect applications from one another
 - Provide efficient access to resources (cost, time, energy)
 - Provide fair access to resources
- **Challenges?**
 - What are the correct mechanisms?
 - What are the correct policies?

Resources are managed via services

- Program Execution (loading, running, monitoring, terminating)
- Performance (optimizing resources under constraints)
- Correctness (overseeing critical operations, preventing interference)
- Fairness (access to and allocation of resources)
- Error detection & recovery (network partition & media failure)

Services (cont'd)

- Communication (inter-process, software-to-hardware, hardware-to-hardware, system-to-system, wide-area)
- I/O: reading & writing, support for various mediums, devices, performance, and protections
- Data Organization (naming), Services (search) & Protection (access control)
- Security (isolation, enforcement, services, authentication, accounting and logging, trust)
- User interfaces (command-line, GUIs, multiple users)

Each service has challenges and tensions

Example 1: We have limited RAM, and we want to run more programs that can be stored.

- How do we allocate space?
- Who stays?
- Who goes?
- What if we're wrong?
- What if the system is under extremely heavy load?
- Is there a way to predict the future?

Each service has challenges and tensions

Example 2: We have two process (producer / consumer); how do they communicate?

- Message passing? Shared memory?
- How do they synchronize?
 - How to we prevent over-production? Over-consumption?
 - Context-switching?

What isn't clear?

Comments? Thoughts?